

CS4223 Assignment 2 Report

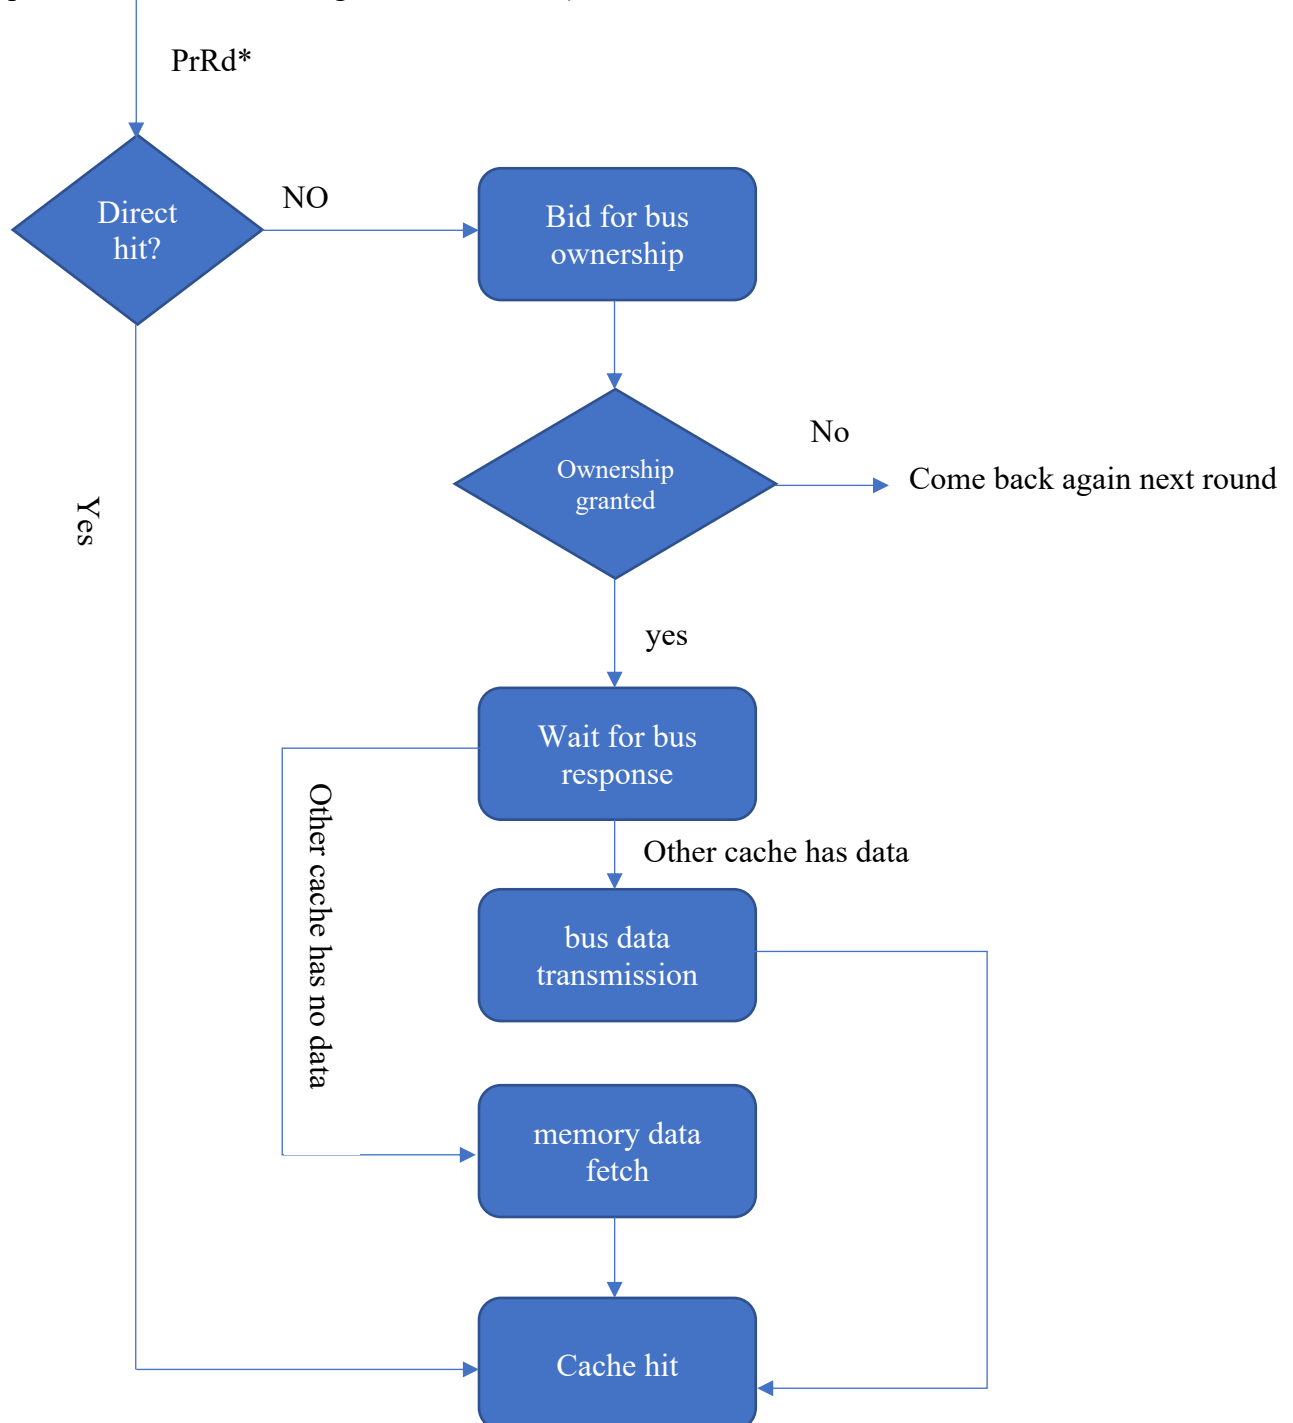
Additional Assumptions

- Memory Controller
 1. Memory can handle infinite concurrent fetches and each fetch takes 100 cycles sharp.
- Bus
 1. At any point of time, only one cache can own the bus. Ownership can only be released by the owner itself.
 2. Bidding and the subsequent granting of bus ownership does not take time, it happens at the instant right before a clock cycle starts, in other words it happens during the *interim* between two clock cycles. Bidding is only allowed if the bus currently has no owner.
 3. Bus ownership can only be granted to one cache once per *interim*. Once a cache becomes the bus owner, other bidders have to back off and wait till the beginning of next cycle to come back and check the availability again, even though the bus ownership may immediately be released by its owner (e.g. a BusRdX that only causes invalidation but no data transfer)
- Locking and Atomicity
 1. In addition to the cache line states, we introduce a 3-state locking bit for each cache line: **unlocked**, **read_locked**, **write_locked**.
 2. A cache block in **unlocked** state is subject to reads and invalidations by all incoming bus requests.
 3. A **read_locked** block can respond to BusRd, but BusRdX or BusUpd requests will stall and only get their response after the lock is released.
 4. A **write_locked** block will stall all incoming bus requests that is targeting that block, only responding to them after the lock is released.
 5. During a read hit, the cache block will be **read_locked**. During a write hit it shall be **write_locked**.
 6. When fetching a block from memory, the block being fetched will be **write_locked**.
 7. When fetching a block from other caches thru the bus, the block will be **write_locked**.
 8. When a bus request for a block does not return positive answers (meaning the block is not found in any of the other caches), the bus ownership will be released. The cache will proceed to fetch the block from memory. Although the bus can now be grabbed by other caches and other cache may want to read this block, due to the aforementioned write lock still in place, access to this block is forbidden thus retaining atomicity for the entire operation.
- Block replacement
 1. When a processor tries to access an address that is not to be found in the corresponding cache set, the eviction of the least recently accessed block **will not happen immediately**. The eviction will only commence when the cache successfully secures the ownership of the bus. If the bus ownership is in other cache's hands, the block will stay for now so whoever owns the bus may still send a BusRd/BusRdX to access/invalidate it before it gets evicted.

2. If any cache access results in a block eviction, the cache will be stalled, waiting for the eviction to complete before proceeding with either bus transaction or memory fetch, effectively adding 100 cycles to the operation.
3. In the case where local Modified block gets flushed by incoming BusRdX, we assume such passive eviction runs synchronously, i.e. eviction will stall the cache.
4. There is no memory bandwidth limit for each core, i.e. multiple memory transactions can be run simultaneously.

Implementation Flowchart

(both protocols share similar high level workflow)



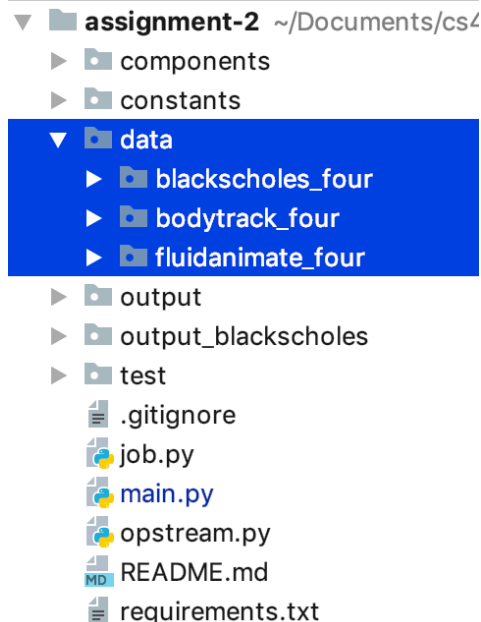
* Flow chart for PrWt is very similar, only that from *wait for bus* you can go directly to *Cache Hit*, if local cache already has data in Shared state.

Key Data Structure

- Cache Data Representation
 - A 3-dimensional numpy array of shape (m, n, 3)
 - m is the number of cache sets given by $\text{cache_size}/\text{associativity}/\text{block_size}$
 - n is the associativity
 - last dimension represents 3-element tuple of [tag, state, lockbit]
- Processor Class
 - In charge of issuing new instruction if available.
 - Counting for non-cache instructions.
- Cache Class
 - Maintains the simulated cache data.
 - Implementation of different state transition mechanisms provided by different protocol.
- Bus Class
 - Connect each cache class.
 - Support bus-initiated data transistions.

Running the Simulator

- Language
 - Python 3
- Setup
 - `pip install -r requirements.txt`
- folder structure
 - make sure the data folder is put under the root directory of the source code.



- Run, e.g.
 - `./coherence.sh MESI ./data/bodytrack_four 1024 1 16`
- Sample Output

```

Run: main x
All Finished! Current counter: 39703006
processor 0 counter: 39272184
processor 1 counter: 39703006
processor 2 counter: 18846743
processor 3 counter: 39283717

stats
Bus Data Traffic      5292992
Bus Invalidation/Updates 59541
Overall Execution Cycle 39703006
Private Data Access Percentage 69

P0      P1      P2      P3
Compute Cycles 17729254.00 17120545.00 17556877.00 17140113.00
Load/Store Instructions 3270132.00 3287252.00 117698.00 3324919.00
Idle Cycles 21542930.00 22582461.00 1289866.00 22143604.00
Cache Miss Rate 0.03 0.04 0.04 0.03

Process finished with exit code 0

```

Advanced Task

Improvement to MESI

MESI-based protocol is more commonly used in modern multi-core architecture. For example, Intel used MESIF for its Core Series CPU. Among many variants of MESI-based protocols, we select MOESI as our additional coherence protocol.

Compared with MESI, MOESI simply add an OWNED state to the MOESI to reduce redundant memory accesses. When a cache line in M state, and a BusRd occurs, this cache line will transition to O state, rather than writing back to memory. O state is similar to S state, except that cache in O state have to write back to main memory, when this cache is to be replaced. The following figure shows the FSM of MOESI.

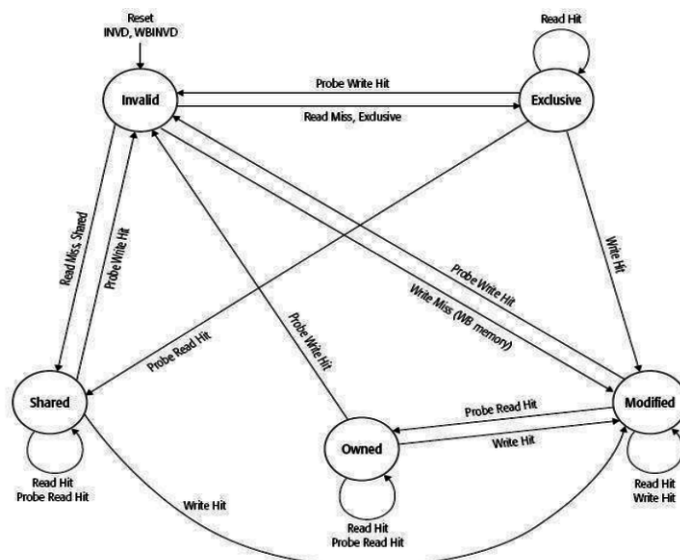


Fig source: Dey, Somdip, and Mamatha S. Nair. "Design and Implementation of a Simple Cache Simulator in Java to Investigate MESI and MOESI Coherency Protocols." *International Journal of Computer Applications* 87.11 (2014).

Improvement to Dragon

For the Dragon protocol, there are some redundant memory operations shown in many state transition graphs. Just like the figure below, if a cache block is in Sm or M state, and a BusRd occurs, then the cache block will be flush to the main memory. Such operation is actually redundant if we consider memory write-back every time caches with Sm or M state is

replaced. In this way, our improvement is to remove the memory flush in Sm and M state. We name this modified protocol Dragon-noflush

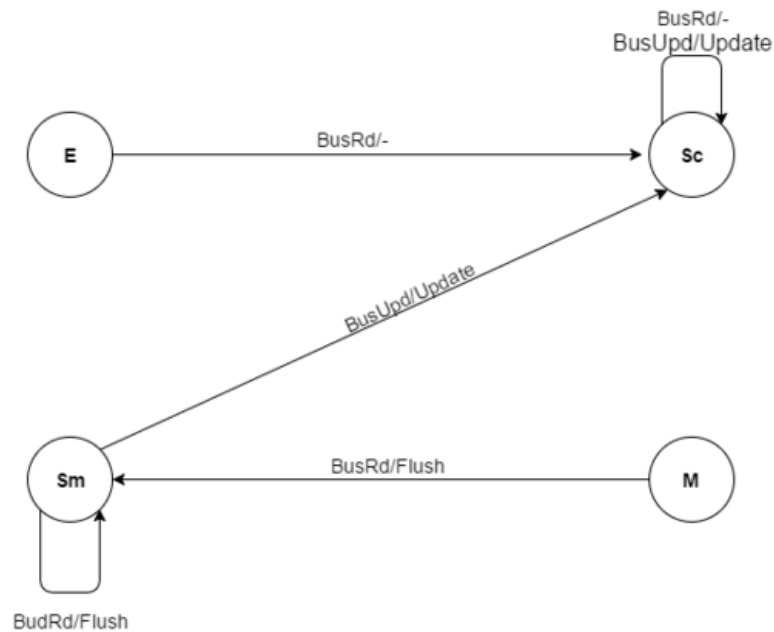


Fig source: wikipedia

Result Analysis

(Configuration: 4KB 2-way associative cache, block size 32 bytes, word size 4 bytes)

1. Black Scholes

| | MESI | | | |
|--------------------------------|----------|----------|----------|----------|
| | P0 | P1 | P2 | P3 |
| Compute Cycles | 10430314 | 10383276 | 10430338 | 10394904 |
| Load/Store Instructions | 2497349 | 2490468 | 2509057 | 2503127 |
| Idle Cycles | 6636123 | 6589273 | 7549819 | 6726709 |
| Cache Miss Rate | 0.01 | 0.01 | 0.01 | 0.01 |
| Bus Data Traffic (Byte) | 697216 | | | |
| Bus Invalidation/Updates | 23290 | | | |
| Overall Execution Cycle | 17980157 | | | |
| Private Data Access Percentage | 81 | | | |
| | DRAGON | | | |
| | P0 | P1 | P2 | P3 |
| Idle Cycles | 6640924 | 6610163 | 7556377 | 6734058 |
| Cache Miss Rate | 0.01 | 0.01 | 0.01 | 0.01 |
| Bus Data Traffic (Byte) | 715456 | | | |
| Bus Invalidation/Updates | 26496 | | | |
| Overall Execution Cycle | 17986715 | | | |
| Private Data Access Percentage | 81 | | | |

| | MOESI | | | |
|--------------------------------|----------------|---------|---------|---------|
| | P0 | P1 | P2 | P3 |
| Idle Cycles | 6635323 | 6575998 | 7547403 | 6725986 |
| Cache Miss Rate | 0.01 | 0.01 | 0.01 | 0.01 |
| Bus Data Traffic (Byte) | 693408 | | | |
| Bus Invalidation/Updates | 23293 | | | |
| Overall Execution Cycle | 17977741 | | | |
| Private Data Access Percentage | 81 | | | |
| | DRAGON-noflush | | | |
| | P0 | P1 | P2 | P3 |
| Idle Cycles | 6643822 | 6603519 | 7556814 | 6732024 |
| Cache Miss Rate | 0.01 | 0.01 | 0.01 | 0.01 |
| Bus Data Traffic (Byte) | 693408 | | | |
| Bus Invalidation/Updates | 23293 | | | |
| Overall Execution Cycle | 17987152 | | | |
| Private Data Access Percentage | 81 | | | |

2. Body Track

| | MESI | | | |
|--------------------------------|----------|----------|----------|----------|
| | P0 | P1 | P2 | P3 |
| Compute Cycles | 17729254 | 17120545 | 17556877 | 17140113 |
| Load/Store Instructions | 3270132 | 3287252 | 117698 | 3324919 |
| Idle Cycles | 21819530 | 22859061 | 1299568 | 22420204 |
| Cache Miss Rate | 0.03 | 0.04 | 0.04 | 0.03 |
| Bus Data Traffic (Byte) | 4984160 | | | |
| Bus Invalidation/Updates | 59476 | | | |
| Overall Execution Cycle | 39979606 | | | |
| Private Data Access Percentage | 71 | | | |
| | DRAGON | | | |
| | P0 | P1 | P2 | P3 |
| Idle Cycles | 21675503 | 22715113 | 1288849 | 22276250 |
| Cache Miss Rate | 0.03 | 0.04 | 0.04 | 0.03 |
| Bus Data Traffic (Byte) | 5064828 | | | |
| Bus Invalidation/Updates | 60975 | | | |
| Overall Execution Cycle | 39835658 | | | |
| Private Data Access Percentage | 70 | | | |
| | MOESI | | | |
| | P0 | P1 | P2 | P3 |
| Idle Cycles | 21542930 | 22582461 | 1289866 | 22143604 |
| Cache Miss Rate | 0.03 | 0.04 | 0.04 | 0.03 |
| Bus Data Traffic | 5292992 | | | |
| Bus Invalidation/Updates | 59541 | | | |

| | | | | |
|--------------------------------|----------------|----------|---------|----------|
| Overall Execution Cycle | 39703006 | | | |
| Private Data Access Percentage | 69 | | | |
| | DRAGON-noflush | | | |
| | P0 | P1 | P2 | P3 |
| Idle Cycles | 21645736 | 22685346 | 1296561 | 22246483 |
| Cache Miss Rate | 0.03 | 0.04 | 0.04 | 0.03 |
| Bus Data Traffic | 5053300 | | | |
| Bus Invalidation/Updates | 60880 | | | |
| Overall Execution Cycle | 39805891 | | | |
| Private Data Access Percentage | 71 | | | |

3. Fluid Animate

| | | | | |
|--------------------------------|----------------|----------|----------|----------|
| | MESI | | | |
| | P0 | P1 | P2 | P3 |
| Compute Cycles | 11337782 | 11290799 | 11337671 | 11301515 |
| Load/Store Instructions | 2576503 | 2407844 | 2604189 | 2411465 |
| Idle Cycles | 33174662 | 31502225 | 37946761 | 38935559 |
| Cache Miss Rate | 0.03 | 0.02 | 0.03 | 0.02 |
| Bus Data Traffic (Byte) | 1168832 | | | |
| Bus Invalidation/Updates | 287074 | | | |
| Overall Execution Cycle | 50237074 | | | |
| Private Data Access Percentage | 84 | | | |
| | DRAGON | | | |
| | P0 | P1 | P2 | P3 |
| Idle Cycles | 33132852 | 31431073 | 37840315 | 38870638 |
| Cache Miss Rate | 0.03 | 0.02 | 0.03 | 0.02 |
| Bus Data Traffic (Byte) | 1119728 | | | |
| Bus Invalidation/Updates | 289340 | | | |
| Overall Execution Cycle | 50172153 | | | |
| Private Data Access Percentage | 85 | | | |
| | MOESI | | | |
| | P0 | P1 | P2 | P3 |
| Idle Cycles | 33077951 | 31380752 | 37798948 | 38787746 |
| Cache Miss Rate | 0.03 | 0.02 | 0.03 | 0.02 |
| Bus Data Traffic (Byte) | 1171584 | | | |
| Bus Invalidation/Updates | 287074 | | | |
| Overall Execution Cycle | 50089261 | | | |
| Private Data Access Percentage | 84 | | | |
| | DRAGON-noflush | | | |
| | P0 | P1 | P2 | P3 |
| Idle Cycles | 33055461 | 31350984 | 37780509 | 38747811 |

| | | | | |
|--------------------------------|----------|------|------|------|
| Cache Miss Rate | 0.03 | 0.02 | 0.03 | 0.02 |
| Bus Data Traffic (Byte) | 1116296 | | | |
| Bus Invalidation/Updates | 289397 | | | |
| Overall Execution Cycle | 50049326 | | | |
| Private Data Access Percentage | 85 | | | |

From the testing results of 3 cache coherence protocols on different benchmark trace, we can see that different protocols perform differently upon the benchmark used. We try to make some analysis based on the results.

1. MESI and DRAGON perform quite similarly on 3 benchmarks under the same cache configuration. For the *blackscholes*, MESI has fewer overall execution cycles, while for *bodytrack* and *fluidanimate*, DRAGON outperforms MESI. To analyze, we know that different coherence protocols can cause different number of accesses to main memory, and different amount of data transferred via bus. Except for *fluidanimate* benchmark, DRAGON has more bus data traffic. By considering the differences between invalidation protocol and update protocol, more data write operations will cause more performance loss for update protocol, which also implies that our testing benchmarks generally have a considerable number of data write operations.
2. MOESI and DRAGON-noflush perform better than their corresponding naïve version in our testing benchmarks. Thanks to the OWNED state, overall execution cycles of MOESI are reduced in all testing benchmarks. For the DRAGON-noflush, the overall execution cycles for *bodytrack* and *fluidanimate* are reduced, while overall execution cycles for *blackscholes* are increased a bit, which may be caused by the different coordination of bus requests.

To conclude, the MOESI protocol has the best overall performance of all three testing benchmarks, which also shows the reason why modern CPU prefer to use MESI or MOESI alike protocols.